

Salary Prediction Analysis

Overview

This repository contains a comprehensive analysis of the "Salary Data.csv" dataset, aiming to predict salaries based on several features like age, years of experience, gender, and job title. Through extensive data processing, visualization, and machine learning, we've derived valuable insights and built predictive models with impressive accuracy.

Dataset

The primary dataset used is "Salary Data.csv", which encompasses several features:

- **Age:** The age of the employees.
- **Years of Experience:** The total working experience of the employees.
- **Gender:** The gender of the employees.
- **Job Title:** The designation or role of the employees.
- **Education Level:** The highest educational qualification of the employees.
- **Salary:** The annual salary of the employees (our target variable).

Procedures and Analysis

Data Exploration and Visualization:

- Loaded the dataset into a pandas DataFrame for easy manipulation and analysis.
- Conducted a comprehensive overview of the dataset, understanding its structure and content.
- Used statistical measures to understand the distribution and central tendencies of the data.
- Visualized the data using libraries like `matplotlib` and `seaborn` to understand the relationships between different features and the target variable.

Data Preprocessing:

- Handled missing values and outliers.
- Encoded categorical variables to convert them into a format suitable for machine learning.
- Split the dataset into training and testing sets to train and subsequently evaluate our machine learning models.

Model Training and Evaluation:

- Employed several regression algorithms, including:

- Linear Regression
 - Decision Tree Regression
 - Extra Tree Regression
 - Gradient Boosting Regression
 - Random Forest Regression
 - Ridge Regression
 - Lasso Regression
 - Elastic Net Regression
- Trained each model on the training set and evaluated their performance on the testing set using R-squared as the metric.

Results

The R-squared values for each of the models are as follows:

- **Linear Regression:** 0.8556
- **Decision Tree Regression:** 0.8889
- **Extra Tree Regression:** 0.8997
- **Gradient Boosting Regression:** 0.9160
- **Random Forest Regression:** 0.9111
- **Ridge Regression:** 0.9036
- **Lasso Regression:** 0.9027
- **Elastic Net Regression:** 0.9046

The Gradient Boosting Regression model showed the highest R-squared value, indicating it performed the best among all the models in terms of explaining the variance in the dataset.

Importing the necessary libraries

```
In [ ]: import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler, OneHotEncoder, OrdinalEncoder
from sklearn.compose import ColumnTransformer, TransformedTargetRegressor
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LinearRegression, Lasso, SGDRegressor, Ridge, ElasticNet
from sklearn.ensemble import GradientBoostingRegressor, RandomForestRegressor
from sklearn.tree import DecisionTreeRegressor, ExtraTreeRegressor
from sklearn.model_selection import train_test_split, RandomizedSearchCV
from sklearn.svm import SVR

plt.style.use('fivethirtyeight')
```

Reading the dataset

```
In [ ]: df = pd.read_csv('Salary Data.csv')
print(df.shape)
df.head()
```

(375, 6)

Out []:

	Age	Gender	Education Level	Job Title	Years of Experience	Salary
0	32.0	Male	Bachelor's	Software Engineer	5.0	90000.0
1	28.0	Female	Master's	Data Analyst	3.0	65000.0
2	45.0	Male	PhD	Senior Manager	15.0	150000.0
3	36.0	Female	Bachelor's	Sales Associate	7.0	60000.0
4	52.0	Male	Master's	Director	20.0	200000.0

An overview of the dataset

In []: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 375 entries, 0 to 374
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age                   373 non-null   float64
1   Gender                373 non-null   object
2   Education Level       373 non-null   object
3   Job Title             373 non-null   object
4   Years of Experience    373 non-null   float64
5   Salary                373 non-null   float64
dtypes: float64(3), object(3)
memory usage: 17.7+ KB
```

Statistical information of numeric features of the dataset

In []: `print(df.describe())`

	Age	Years of Experience	Salary
count	373.000000	373.000000	373.000000
mean	37.431635	10.030831	100577.345845
std	7.069073	6.557007	48240.013482
min	23.000000	0.000000	350.000000
25%	31.000000	4.000000	55000.000000
50%	36.000000	9.000000	95000.000000
75%	44.000000	15.000000	140000.000000
max	53.000000	25.000000	250000.000000

Count of unique values for each categorical columns


In []: `for col in ('Gender', 'Education Level', 'Job Title'):`
`print(f'for column {col}, there is {df[col].nunique()} unique values.')`

```
for column Gender, there is 2 unique values.
for column Education Level, there is 3 unique values.
for column Job Title, there is 174 unique values.
```

EDA (Exploratory data analysis)

Scatter plot for Salary and Age features

```
In [ ]: plt.scatter(df['Age'], df['Salary'])  
plt.xlabel('Age')  
plt.ylabel('Salary')  
plt.title('Scatter plot for Salary and Age')  
plt.show()
```




Scatter plot for Salary and Years of Experience

```
In [ ]: plt.scatter(df['Years of Experience'], df['Salary'])  
plt.xlabel('Years of Experience')  
plt.ylabel('Salary')  
plt.title('Scatter plot for Salary and Years of Experience')  
plt.show()
```



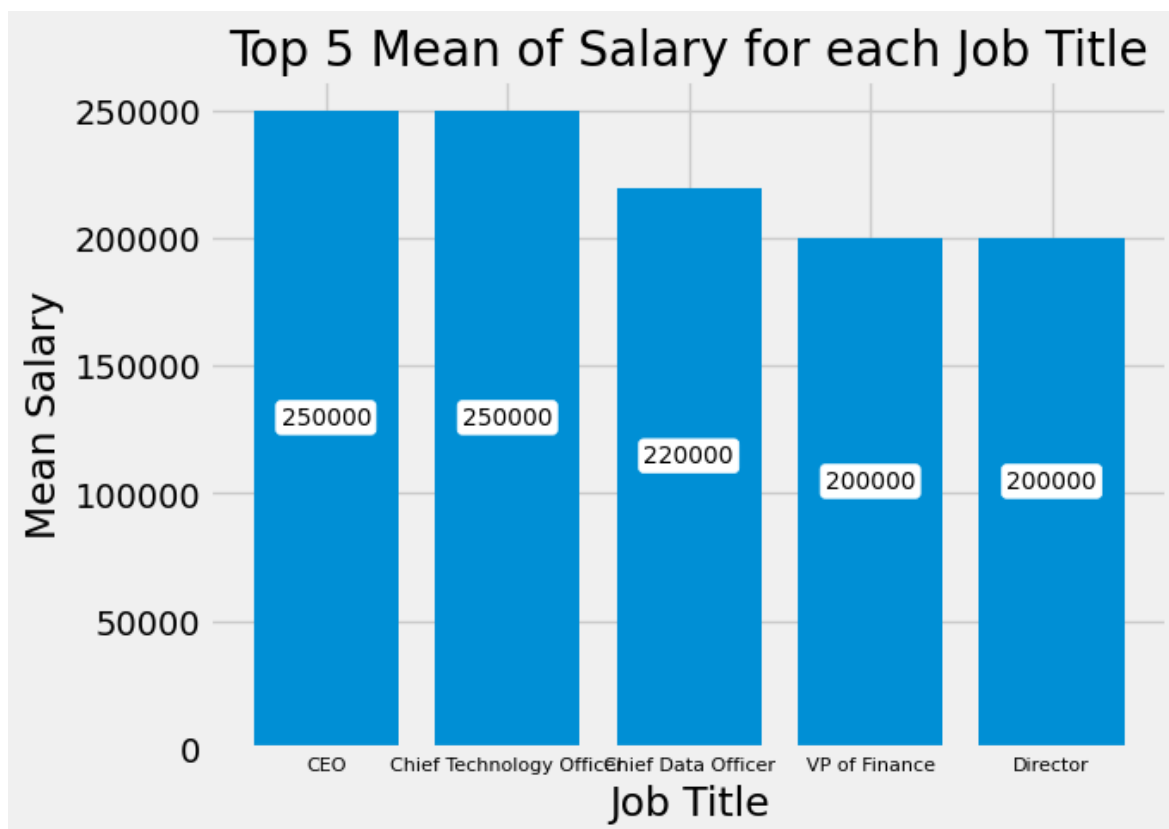
Mean of Salary for each Educational Level

```
In [ ]: educationSalary = df.groupby('Education Level').agg({'Salary': lambda x: x.mean()})
bars = plt.bar(educationSalary['Education Level'], educationSalary['Salary'])
for bar in bars:
    plt.text(bar.get_x() + bar.get_width()/2, bar.get_height()/2, int(bar.get_h
plt.title('Mean of Salary for each Educational Level')
plt.xlabel('Education Level')
plt.ylabel('Mean Salary')
plt.show()
```



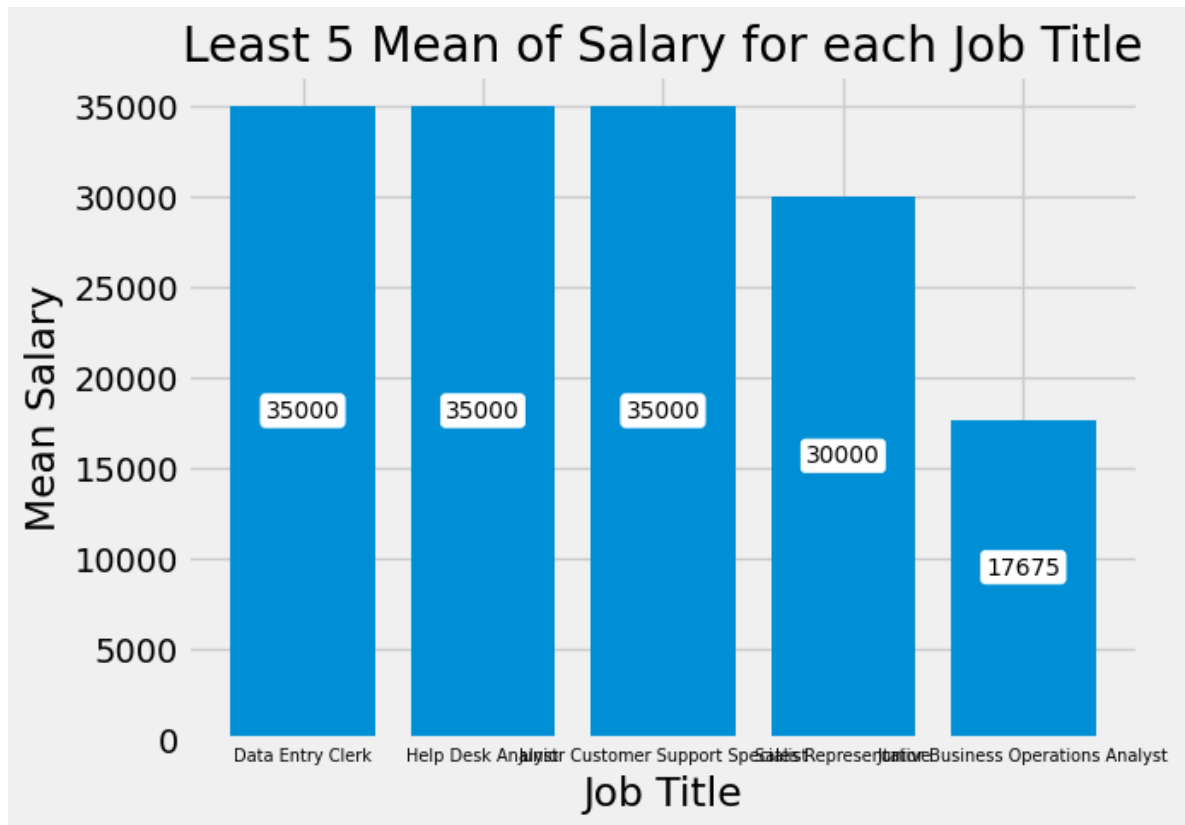
Top 5 Mean of Salary for each Job Title

```
In [ ]: jobTitleSalary = df.groupby('Job Title').agg({'Salary': lambda x: x.mean()}).sort_values(
bars = plt.bar(jobTitleSalary['Job Title'], jobTitleSalary['Salary'])
for bar in bars:
    plt.text(bar.get_x() + bar.get_width()/2, bar.get_height()/2, int(bar.get_h
plt.xticks(fontsize=8)
plt.title('Top 5 Mean of Salary for each Job Title')
plt.xlabel('Job Title')
plt.ylabel('Mean Salary')
plt.show()
```




Least 5 Mean of Salary for each Job Title

```
In [ ]: jobTitleSalary = df.groupby('Job Title').agg({'Salary': lambda x: x.mean()}).sort_values(
        bars = plt.bar(jobTitleSalary['Job Title'], jobTitleSalary['Salary'])
        for bar in bars:
            plt.text(bar.get_x() + bar.get_width()/2, bar.get_height()/2, int(bar.get_h
        plt.xticks(fontsize=7)
        plt.title('Least 5 Mean of Salary for each Job Title')
        plt.xlabel('Job Title')
        plt.ylabel('Mean Salary')
        plt.show()
```



Mean of Salary for each Gender

```
In [ ]: genderSalary = df.groupby('Gender').agg({'Salary': lambda x: x.mean()}).reset_index()
bars = plt.bar(genderSalary['Gender'], genderSalary['Salary'])
for bar in bars:
    plt.text(bar.get_x() + bar.get_width()/2, bar.get_height()/2, int(bar.get_height()))
plt.xticks(fontsize=10)
plt.title('Mean of Salary for each Gender')
plt.xlabel('Gender')
plt.ylabel('Mean Salary')
plt.show()
```

Defining our data workflow

Preprocessing our data using ColumnTransformer

```
In [ ]: import numpy as np

cat_features = ['Gender', 'Education Level', 'Job Title']
num_features = ['Age', 'Years of Experience']

df.dropna(inplace=True)

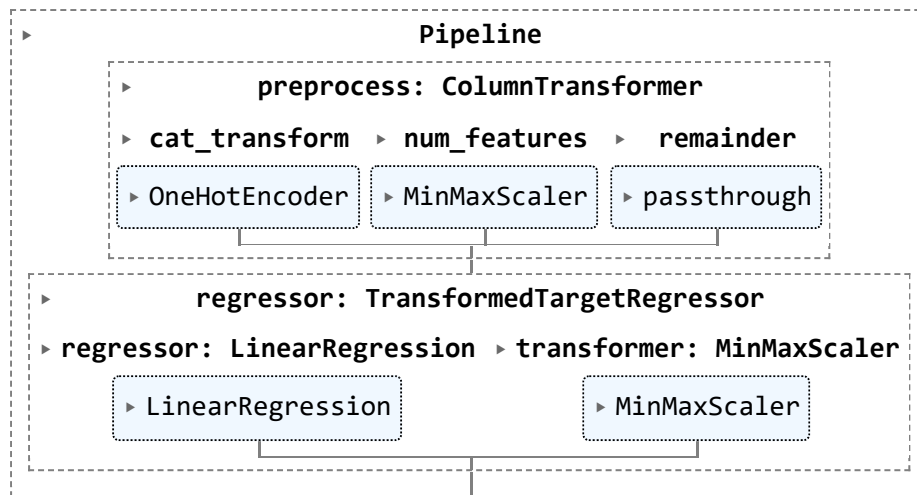
preprocess = ColumnTransformer([
    # ('cat_transform', OrdinalEncoder(handle_unknown='use_encoded_value', unknown_value=-1), df.columns.get_indexer_for(cat_features)),
    ('cat_transform', OneHotEncoder(handle_unknown='infrequent_if_exist'), df.columns.get_indexer_for(cat_features)),
    ('num_features', MinMaxScaler(), df.columns.get_indexer_for(num_features)),
], remainder='passthrough')
```

Enhancing our data processing using Pipeline

```
In [ ]: pipeline = Pipeline([
    ('preprocess', preprocess),
    ('regressor', TransformedTargetRegressor(regressor=LinearRegression(), transformer=None))
])

X, y = df.iloc[:, :-1], df.iloc[:, -1]
X_train, X_test, y_train, y_test = train_test_split(X.values, y.values, train_size=0.8, random_state=42)
pipeline.fit(X_train, y_train)
```

Out[]:



Defining a dictionary for hyperparameter distributions

```

In [ ]: param_distributions = {
    'LinearRegression': {
        'regressor__fit_intercept': [True, False],
    },
    'DecisionTreeRegressor': {
        'regressor__max_depth': [None] + list(np.arange(1, 21)),
        'regressor__min_samples_split': np.arange(2, 21),
        'regressor__min_samples_leaf': np.arange(1, 21)
    },
    'ExtraTreeRegressor': {
        'regressor__max_depth': [None] + list(np.arange(1, 21)),
        'regressor__min_samples_split': np.arange(2, 21),
        'regressor__min_samples_leaf': np.arange(1, 21),
    },
    'GradientBoostingRegressor': {
        'regressor__n_estimators': np.arange(50, 501, 50),
        'regressor__learning_rate': [0.001, 0.01, 0.05, 0.1, 0.2],
        'regressor__max_depth': np.arange(1, 11),
        'regressor__min_samples_split': np.arange(2, 21),
        'regressor__min_samples_leaf': np.arange(1, 21),
        'regressor__subsample': np.linspace(0.5, 1, 6)
    },
    'RandomForestRegressor': {
        'regressor__n_estimators': np.arange(50, 501, 50),
        'regressor__max_depth': [None] + list(np.arange(1, 21)),
        'regressor__min_samples_split': np.arange(2, 21),
        'regressor__min_samples_leaf': np.arange(1, 21),
        'regressor__bootstrap': [True, False]
    },
    'Ridge': {
        'regressor__alpha': np.logspace(-6, 6, 13),
        'regressor__fit_intercept': [True, False],
    },
    'Lasso': {
        'regressor__alpha': np.logspace(-6, 6, 13),
        'regressor__fit_intercept': [True, False],
        # 'regressor__normalize': [True, False],
        'regressor__selection': ['cyclic', 'random']
    },
}

```

```

'ElasticNet': {
    'regressor__alpha': np.logspace(-6, 6, 13),
    'regressor__l1_ratio': np.linspace(0, 1, 11),
    'regressor__fit_intercept': [True, False],
    'regressor__selection': ['cyclic', 'random']
}

```

Hyperparameter optimizing using RandomizeSearchCV

```

In [ ]: best_models = dict()
best_params = dict()
best_score = dict()
X_train, X_test, y_train, y_test = train_test_split(X.values, y.values, train_si
for model in param_distributions.keys():
    print(model)
    pipeline.steps[1] = (f"regressor", eval(f"{model}()"))
    random_search = RandomizedSearchCV(pipeline, param_distributions=param_distr
                                   n_iter=50, cv=5, random_state=42)

    random_search.fit(X_train, y_train)
    best_models[model] = random_search.best_estimator_
    best_params[model] = random_search.best_params_
    best_score[model] = random_search.best_score_

```


Plotting the models performances

- we can see that the GradientBoostingRegressor has the best performance

```

In [ ]: fig = plt.figure(figsize=(12, 5))
df_best_models = pd.DataFrame().from_dict(best_score, orient='index').reset_index()
df_best_models.columns = ['model', 'score']
plt.plot(df_best_models['model'], df_best_models['score'], marker='o', ms=10, li
plt.xticks(fontsize=10, rotation=30)
plt.xlabel('Models')
plt.ylabel('Score')
plt.show()


```



Plotting the scatter chart of predicted results w.r.t true target variable

```
In [ ]: model = 'GradientBoostingRegressor'
y_pred = best_models[model].predict(X_test)
# y_pred = pipeline.predict(X_test)

plt.figure(figsize=(10, 7))
plt.scatter(y_test, y_pred, alpha=0.6)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'k--', lw=2)
plt.xlabel('True Salary')
plt.ylabel('Predicted Salary')
plt.title(f'True vs. Predicted Salary ({model} Regressor)')
plt.grid(True)
plt.show()
```



```
In [ ]: best_score
```

```
Out[ ]: {'LinearRegression': 0.8556488025821867,  
        'DecisionTreeRegressor': 0.8889224873500357,  
        'ExtraTreeRegressor': 0.8997287368978805,  
        'GradientBoostingRegressor': 0.9160396975265999,  
        'RandomForestRegressor': 0.9111265431473127,  
        'Ridge': 0.903617312335013,  
        'Lasso': 0.9027462613325834,  
        'ElasticNet': 0.9046184977910675}
```

```
In [ ]:
```